

Neural Networks

An Artificial Neural Network consists of an input layer, a number of hidden layers and finally an output layer.

Assume the data are vectors in \mathbb{R}^d i.e. d -dimensional tuples of real numbers.

We are counting the layers from the top so the output layer is layer 0 and if there are a total of N layers, the input layer is layer N , so the hidden layers are numbered $1, 2, \dots, N - 1$

Each layer consists of 3 pieces of data

- A matrix W
- A bias vector \underline{b}
- An activation function $\mathbb{R} \rightarrow \mathbb{R}$, which can be for instance the logistic function σ or the hyperbolic tangent \tanh . Nowadays the most common activation function is the *ReLU*, $x \mapsto \max(x, 0)$

A data vector \underline{x} travels through the network.

$$\underline{x} \mapsto (\underline{u}_N = (\underline{x}W_N + \underline{b}_N)) \mapsto (\phi_N(\underline{u}_N) = \underline{z}_{N-1}) \mapsto (\underline{u}_{N-1} = (\underline{z}_{N-1}W_{N-1} + \underline{b}_{N-1})) \mapsto \dots$$

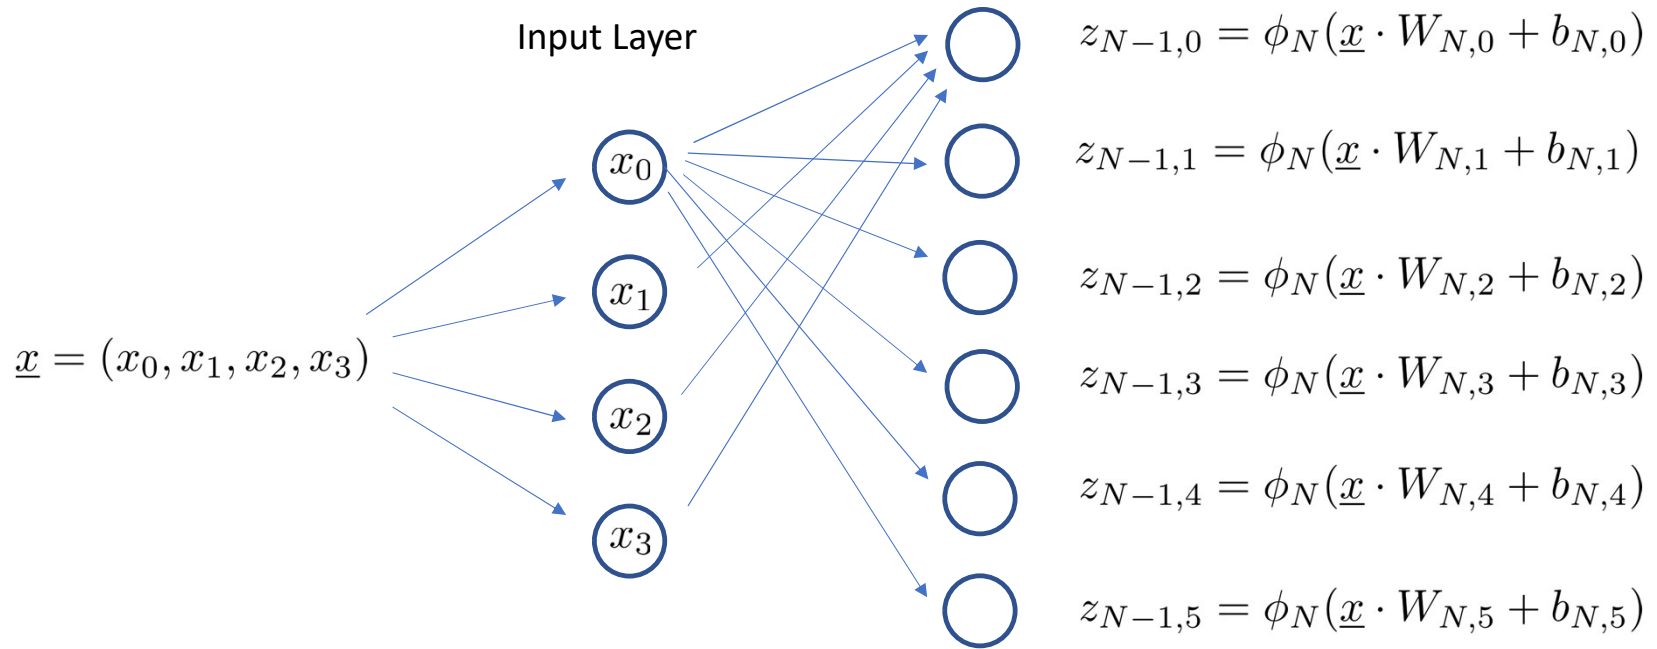
where the matrix W_i , the bias vector \underline{b}_i and the activation function ϕ_i are the data belonging to layer i .

Remark that if W_i is a $d_i \times e_i$ dimensional matrix then \underline{b}_i is an e_i -dimensional vector and W_{i-1} must have e_i rows i.e. W_{i-1} is $d_{i-1} \times e_{i-1}$ with $d_{i-1} = e_i$

Hidden Layer N-1


$W_{N-1} = (W_{N-1,0} \ W_{N-1,1} \ W_{N-1,2} \ W_{N-1,3} \ W_{N-1,4} \ W_{N-1,5})$, of dimension 4×6 so each $W_{N-1,i}$ is a column vector of dimension 4.

Bias vector $\underline{b}_{N-1} = (b_{N-1,0}, b_{N-1,1}, \dots, b_{N-1,5})$ and activation function ϕ_{N-1}



$$\underline{z}_{N-1} = \phi_N(\underline{x}W_N + \underline{b}_N)$$

$$\underline{u}_N = \underline{x}W_N + \underline{b}_N$$



Final activation function , ϕ_1 is typically the softmax or the logistic function (classification) or the identity (regression)

$$\phi_1(\underline{z}_1 W_1 + \underline{b}_1) = \underline{z}_0$$

In a sense the graphic depiction of the ANN is more confusing than illuminating. The reality is that a (in this case *fully connected, feed forward Neural Network*) is just a composite function:

$$f : \underline{x} \mapsto \phi_1(\dots (\phi_{N-1} ((\phi_N (\underline{x}W_N + \underline{b}_N)) W_{N-1} + \underline{b}_{N-1}) W_{N-2} + \underline{b}_{N-2}) \dots$$

or using the z and u variables

$$\begin{aligned} x &\mapsto \underline{z}_0 \\ \underline{z}_0 &= \phi_1(\underline{u}_1) \\ \underline{u}_1 &= \underline{z}_1 W_1 + \underline{b}_1 \\ \underline{z}_1 &= \phi_2(\underline{u}_2) \\ \underline{u}_2 &= \underline{z}_2 W_2 + \underline{b}_2 \\ &\vdots \end{aligned}$$

How do we train an ANN from a dataset i.e. how do we find parameters to fit the data set?

As usual we try to estimate the parameters to minimize a *loss function* \mathcal{L} , which can be for instance Cross Entropy in the classification case or Mean Squared Error (MSE) in the regression case.

The parameters of the ANN consists of the matrices W_N, W_{N-1}, \dots, W_1 and the bias vectors $\underline{b}_N, \underline{b}_{N-1}, \dots, \underline{b}_1$. Thus the total number of parameters is

$$\sum_{i=1}^N d_i e_i + \sum_{i=1}^N e_i = \sum_{i=1}^N d_{i-1} (d_i + 1)$$

Here d_N is the input dimension and d_0 is the output dimension.

For example if $N = 1$ and $\phi_1 = \sigma$, then the ANN is just the logistic regression $\underline{x} \mapsto \sigma(\underline{x}W + b)$, and there are $d_N + 1$ parameters where $d_N = \dim \underline{x}$.

Now assume we have a data point (y, \underline{x}) and assume we have a some loss function \mathcal{L} . We then want to compute $\frac{\partial \mathcal{L}}{\partial W_i}$ for $i = 1, 2, \dots, N$. Using the chain rule we get

$$\frac{\partial \mathcal{L}}{\partial W_i} = \frac{\partial \mathcal{L}}{\partial z_0} \cdot \frac{\partial g_1}{\partial z_1} \cdot \frac{\partial g_2}{\partial z_2} \dots \frac{\partial g_{i-1}}{\partial z_{i-1}} \cdot \frac{\partial g_i}{\partial W_i} = \frac{\partial \mathcal{L}}{\partial z_0} \cdot \nabla \phi_1 \cdot W_1^T \cdot \nabla \phi_2 \cdot W_2^T \dots \nabla \phi_{i-1} \cdot W_{i-1}^T \cdot \frac{\partial g_i}{\partial W_i}$$

The gradients with respect to the the bias vectors are computed similary:


$$\frac{\partial \mathcal{L}}{\partial \underline{b}_i} = \frac{\partial \mathcal{L}}{\partial z_0} \cdot \frac{\partial g_1}{\partial z_1} \cdot \frac{\partial g_2}{\partial z_2} \dots \frac{\partial g_{i-1}}{\partial z_{i-1}} \cdot \frac{\partial g_i}{\partial \underline{b}_i} = \frac{\partial \mathcal{L}}{\partial z_0} \cdot \nabla \phi_1 \cdot W_1^T \cdot \nabla \phi_2 \cdot W_2^T \dots \nabla \phi_{i-1} \cdot W_{i-1}^T \cdot \frac{\partial g_i}{\partial \underline{b}_i}$$

Now we will build a simple Neural Net to perform a classification task.

The dataset is generated by one of the dataset generators in *sklearn*

```
1 X,y = make_moons(500,noise=0.1)
```

```
1 plt.scatter(X[:,0],X[:,1],c=y);
```



We shall use Pytorch to code our deep learning networks (another option is to use TensorFlow but personally I prefer Pytorch)

You need to install pytorch. Go to the Pytorch website to get instructions on how to do this

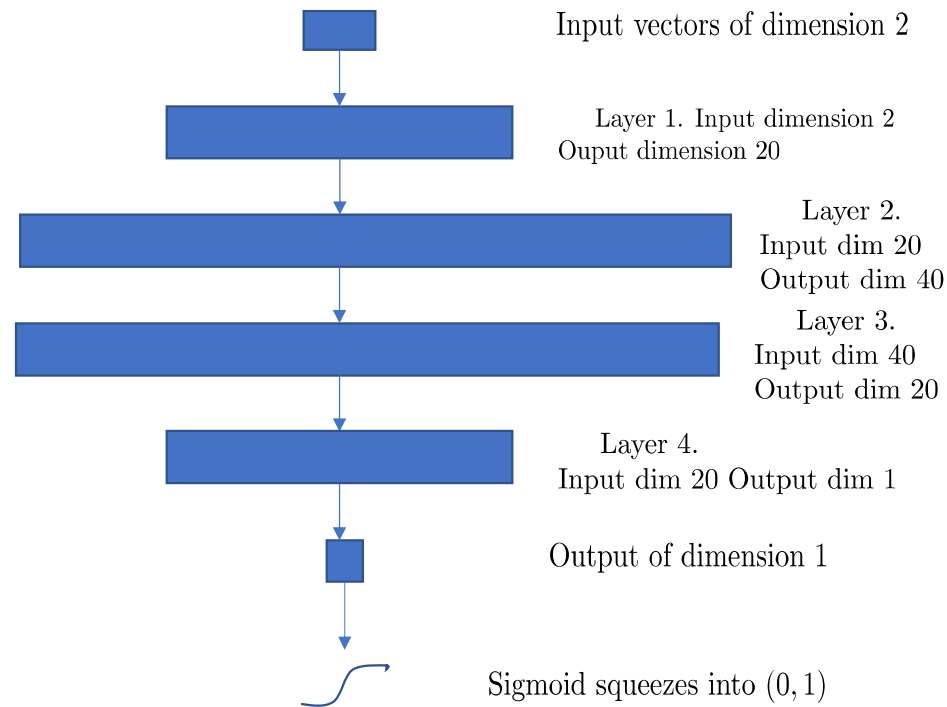
```
1 import numpy as np
2 from sklearn.datasets import make_moons
3 import matplotlib.pyplot as plt
4 %matplotlib inline
```

```
1 import torch
2 import torch.nn as nn
3 import torch.nn.functional as F
4 from torch.optim import Adam
```

Pytorch imports,
remark that we are importing *torch*

We are going to build a network that has an input and an output layer and 4 hidden layers. At the end we send the output through a sigmoid function

$$x \rightarrow \frac{1}{1 + \exp(-x)}$$



```
class Net(nn.Module):
```

In pytorch a network is a class derived from the *nn.Module* class

```
def __init__(self):
```

```
    super().__init__()
```

```
    self.fc1 = nn.Linear(2,20)
```

```
    self.fc2 = nn.Linear(20,40)
```

```
    self.fc3 = nn.Linear(40,20)
```

```
    self.fc4 = nn.Linear(20,1)
```

```
    self.relu = nn.ReLU()
```

Constructor begins by initializing the base class. Then we construct 4 linear layers, remark that the output dimension from a layer has to match the input dimension of the following layer.

The *forward* method overrides the *forward* method in the base class. It sends an input through the entire network. Remark that we have *relu* activations between the linear layers.

```
def forward(self,x):
```

2×20

```
    u = self.fc1(x)     $u = x \cdot W_1 + b_1$ 
```

```
    z = self.relu(u)
```

```
    u = self.fc2(z)
```

```
    z = self.relu(u)
```

```
    u = self.fc3(z)
```

```
    z = self.relu(u)
```

```
    u = self.fc4(z)
```

```
    z = torch.sigmoid(u)
```

The activation for the last layer is the sigmoid function

```
    return z
```

```
def predict(self,A):
```

```
    with torch.no_grad():
```

```
        output = self.forward(A)
```

```
        output.cpu()
```

```
        output_ = output.cpu().numpy()
```

```
        return np.array([0 if x<1/2 else 1 for x in output_])
```

Pytorch (and TensorFlow) deals with *tensors*. Tensors can be viewed simply as multi-dimensional arrays, so for instance a 1-tensor is a vector, a 2-tensor is a matrix (=an array of vectors), a 3-tensor, an array of matrices etc. Pytorch has a special Tensor class and methods that transform arrays into Pytorch tensors.

```
1 x_t = torch.FloatTensor(X)
2 y_t = torch.FloatTensor(y).unsqueeze(1)
```

```
1 x_t.size()
```

```
torch.Size([500, 2])
```

```
1 y_t.size()
```

```
torch.Size([500, 1])
```

```
1 y_t
```

```
[0.],
[0.],
[1.],
[0.],
[0.],
[1.],
[0.],
[1.],
[0.],
[1.],
...
```

The *torch.FloatTensor* turns the arrays into tensors with 32-bit float entries. Pytorch is very picky about tensors having the correct dimensions. Here we need the *y_t* tensor to be a column vector i.e. have column dimension = 1.

We use the *unsqueeze* command to add a dimension.

Next we have to code the *training loop*.

- we feed mini-batches of data points into the network
- we compute the loss between the output from the network and the true labels
- compute the gradient of the loss function on the mini-batch
- the optimizer then does the gradient descent step

Pytorch has a class which will randomly select mini-batches which can be fed into the network, *DataLoader*.

The input to the DataLoader is a *TensorDataset*

```
1 from torch.utils.data import TensorDataset, DataLoader    Construct TensorDataset from the datapoint and the label tensors.
2                                                         Construct a DataLoader that makes mini-batches of length 10
3 dataset = TensorDataset(X_t,y_t)                        and shuffles the data so it comes out in random order
4 dl = DataLoader(dataset = dataset,batch_size=10,
5                 shuffle=True)
```

Next we have to instantiate the *Net* class to get an actual network.

We then specify the loss function, in this case we use Binary Cross Entropy (*nn.BCELoss*). The loss between an output from the network x (which is a number between 0 and 1) and the label y (which is 0 or 1) given by

$$loss_{fn}(x, y) = y \log(x) + (1 - y) \log(1 - x)$$

The optimizer is the Adam optimizer, you should basically always use this optimizer.

```
1 ann = Net()
```

```
1 loss_fn = nn.BCELoss()
```

```
2 optimizer = Adam(ann.parameters(), lr = 0.01)
```

The optimizer needs to know which parameters to work with so it takes as input the parameters in the network.

We also specify a learning rate.

Now we can code the training loop

```
1 losses = []
2
3 for epoch in range(100):
4     aggr_loss = 0.0
5     for points, labels in dl:
6
7         output = ann(points)
8
9         loss = loss_fn(output, labels)
10
11         aggr_loss += loss
12
13         optimizer.zero_grad()
14
15         loss.backward()
16
17         optimizer.step()
18
19 print(aggr_loss.item())
20 losses.append(aggr_loss.item())
```

We want to record the losses from the training

An epoch is one pass through the entire dataset so here we run through the entire dataset (as mini-batches of size 10)

The data loader outputs datapoints and corresponding labels

We send the data points through the network, the output is a tensor of dimension (10,1) with entries, numbers between 0 and 1.

The loss_fn compute the loss between the output and the labels using the BCELoss.

We add the loss to the total loss for the epoch


We set all the gradients = 0

The command *.backward*s automaticall computes all the gradients for all the parameters in the network. This is really the main reason for using a framework like Pytorch or TensorFlow, the auto gradient capability

The *optimizer.step* command then does the gradient descent

The losses show that we probably did not have to train for 100 epochs

```
1 plt.plot(losses);
```



The trained network defines a function $\mathbb{R}^2 \rightarrow 0, 1$

